



Computer Vision

Chapter 3: Image Processing

Computer Vision

Chapter 3. Image Processing

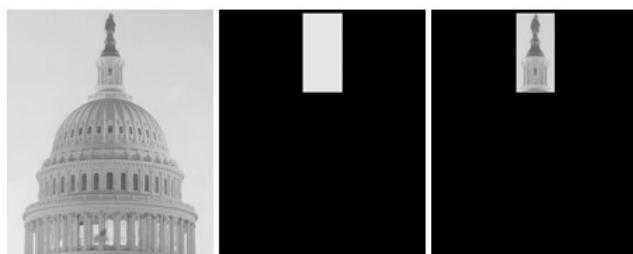
Content

- Rappel: digital image representation
- Point Processing
- Convolution and Linear filtering
- More neighborhood operators
 - Median/max/min filters
 - Arithmetical/Logical operations
 - Binary image and morphological operations
- Image transforms

Arithmetical/Logical Operations

- AND operation
- OR operation
- Image subtraction
- Image addition

AND operation

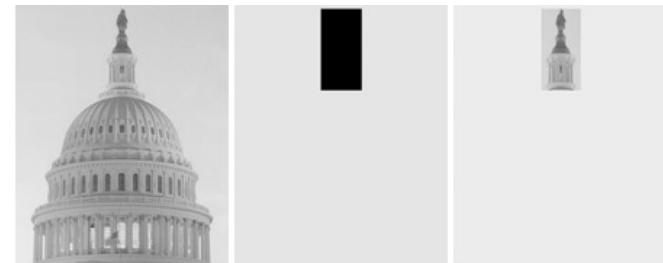


Original

And mask

Output image

OR operation



Original

OR mask

Output image



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

5



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

6

Image Addition

- If f and g are two images, the pixelwise addition R is defined as:

$$- R(x,y) = \text{Min}(f(x,y) + g(x,y) ; 255)$$

- Image addition is used to
 - lower the noise in a serie of images
 - increase the luminance by adding the image to itself



Source : Eric Favier. L'analyse et le traitement des images. ENISE.

7

Average Images

- $g(x,y)$ is the addition of $f(x,y)$ and noise $\eta(x,y)$

$$g(x, y) = f(x, y) + \eta(x, y)$$

- If we have several images $\{g(x,y)\}$, we can compute the average one

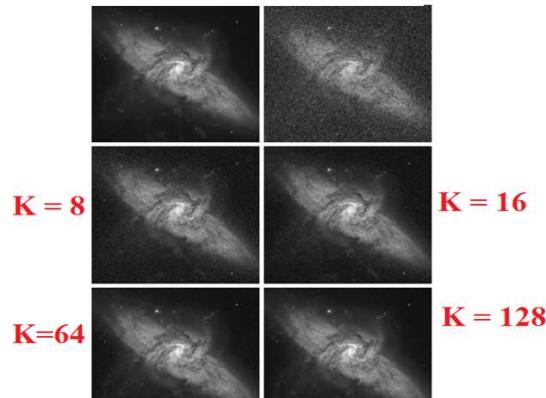
$$\bar{g}(x, y) = \frac{1}{K} \sum_{i=1}^K g_i(x, y)$$



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

8

Average Images



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

9

Image subtraction

- The pixelwise subtraction of two images f and g is:

$$S(x,y) = \text{Max}(f(x,y)-g(x,y) ; 0)$$

- Image subtraction is used to

- detect defaults, detect difference between images
- detect motion in images

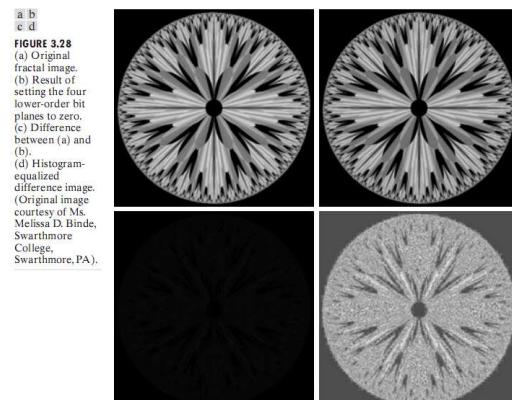


SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Source : Eric Favier. *L'analyse et le traitement des images*. ENISE.

10

Image subtraction



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

11

Image subtraction



After detection, we still have some noise, that we can clean to keep only the object of interest

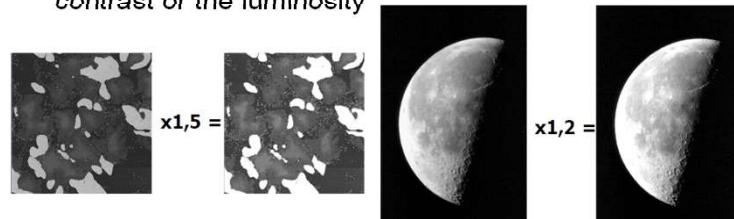


SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

12

Image multiplication

- The multiplication S of an image f by a ratio (factor) is defined as:
 - $S(x,y) = \text{Max}(f(x,y)*\text{ratio} ; 255)$
- Image multiplication can be used to increase the contrast or the luminosity



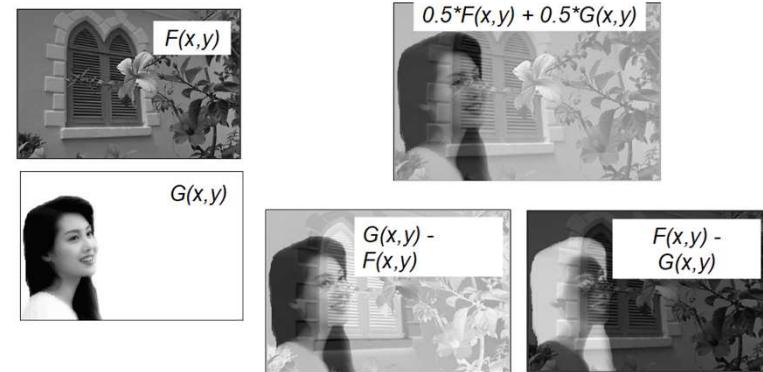
Source : Eric Favier. *L'analyse et le traitement des images*. ENISE.

13



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Operations on images

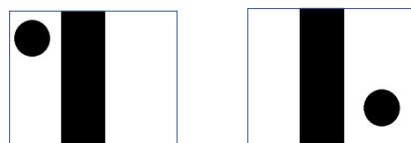


Source : www.nte.montaigne.u-bordeaux.fr/SuppCours/5314/Dai/TraitImage01-02.ppt

14

Exercise

- Given two images as bellow



- Transform images to negative ones
- Process to have an image which has only the "ball"



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

15

Reference

```

• I1 = imread('images/ball1.png');
• >> I1 = rgb2gray(I1);
• >> I1 = 255 - I1;
• >> I2 = imread('images/ball2.png');
• >> I2 = rgb2gray(I2);
• >> I2 = 255 - I2;
• >> I3 = I1 + I2;
• >> I4 = image_subtract(I1, I2);
• >> I5 = I4;
• >> I5(find(I4==255))=255;
• >> I5 = uint8(I5);
• I6 = I3 - I5;
• I7 = I1 - I6;
• I8 = I2 - I6;
  
```



16

Binary images

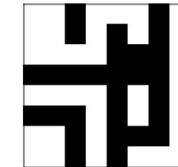


25 SOICT SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Binary images

- Two pixel values: **foreground (object, 1)** and **background (0)**
- Be used
 - To mark region(s) of interest
 - As results of thresholding method

1	1	0	1	1	1	0	1
1	1	0	1	0	1	0	1
1	1	1	1	0	0	0	1
0	0	0	0	0	0	0	1
1	1	1	1	0	1	0	1
0	0	0	1	0	1	0	1
1	1	0	1	0	0	0	1
1	1	0	1	0	1	1	1

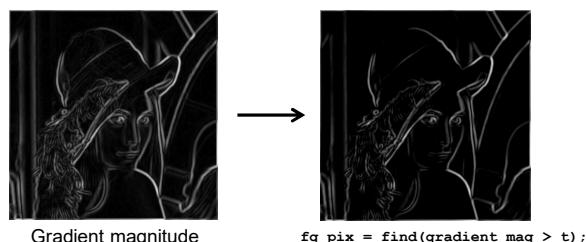


25 SOICT SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Thresholding

- Given a grayscale image or an intermediate matrix → threshold to create a binary output.

Example: edge detection



Gradient magnitude

`fg_pix = find(gradient_mag > t);`

Looking for pixels where gradient is strong.



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Thresholding

- Given a grayscale image or an intermediate matrix → threshold to create a binary output.

Example: background subtraction



Looking for pixels that differ significantly from the "empty" background.

`fg_pix = find(diff > t);`

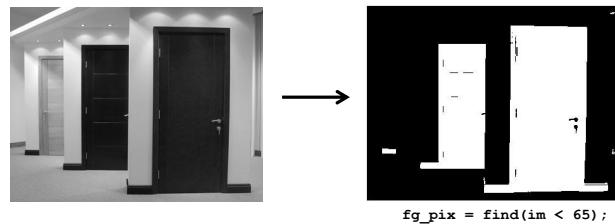


25 SOICT SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Thresholding

- Given a grayscale image or an intermediate matrix → threshold to create a binary output.

Example: intensity-based detection



Looking for dark pixels

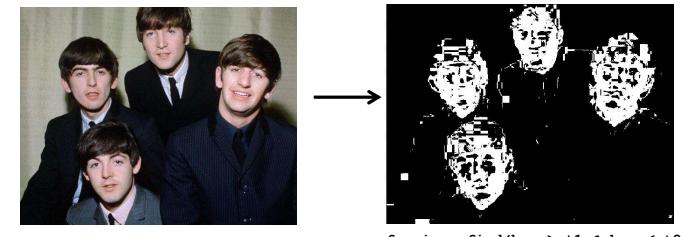


SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Thresholding

- Given a grayscale image or an intermediate matrix → threshold to create a binary output.

Example: color-based detection



Looking for pixels within a certain hue range.

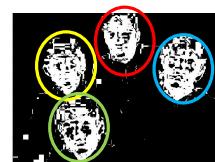


SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Slide credit: Kristen Grauman

Issues

- What to do with “noisy” binary outputs?
 - Holes
 - Extra small fragments
- How to demarcate multiple regions of interest?
 - Count objects
 - Compute further features per object



Morphological operators

- Change the shape of the foreground regions via intersection/union operations between a scanning structuring element and a binary image.
- Useful to clean up result from thresholding
- Main components
 - Structuring element
 - Operators:
 - Basic operators: Dilation, Erosion
 - Others: Opening, Closing, ...



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

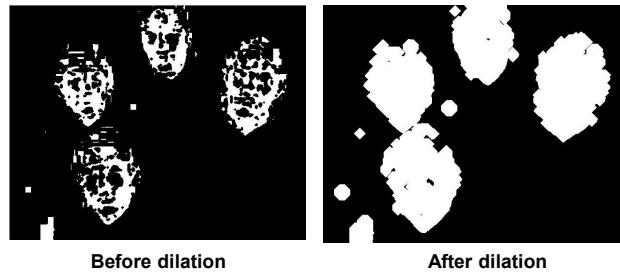
Slide credit: Kristen Grauman



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Dilation

- Expands connected components
- Grow features
- Fill holes



Before dilation

After dilation

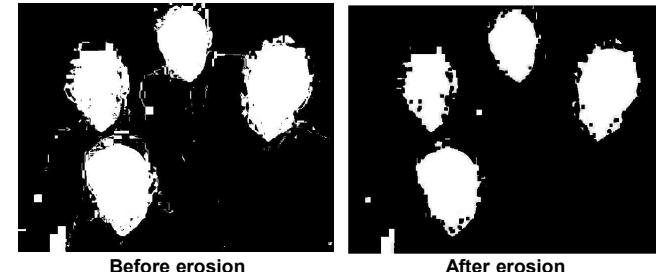
Slide credit: Kristen Grauman



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Erosion

- Erode connected components
- Shrink features
- Remove bridges, branches, noise



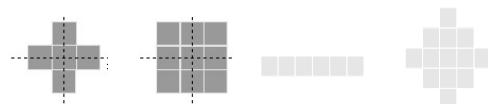
Before erosion

After erosion

Slide credit: Kristen Grauman

Structuring elements

- **Masks** of varying shapes and sizes used to perform morphology, for example:



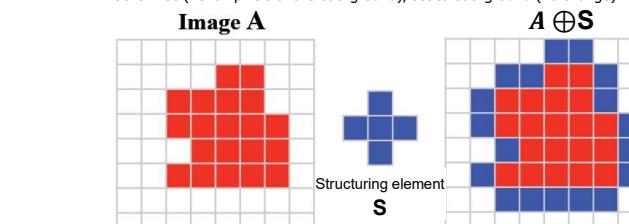
- Scan mask (structuring element) over the **object (foreground) borders (inside and outside)** and transform the binary image



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Dilation

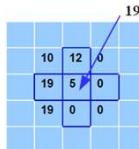
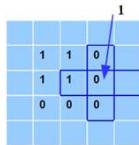
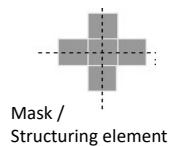
- Moving S on each pixel of A
 - check if the intersection (pixels belonging to object) is not empty
 - If yes, the center of B belongs to the result image
- If a pixel of S is onto object pixels (A), then the **central pixel** belongs to object
 - Otherwise (i.e. all pixels of are background), set to background (no change)

Image A $A \oplus S$
Structuring element S

Structuring element S

Dilation

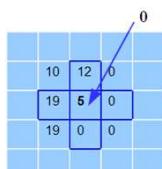
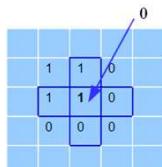
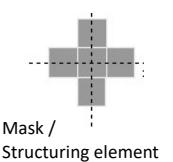
- As **max filter**
- Can be applied both on
 - binary images
 - or grayscale images



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Erosion

- As **min filter**
- Can be applied both on
 - binary images
 - or grayscale images

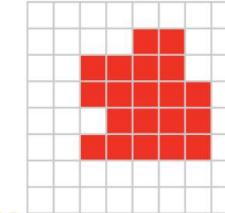
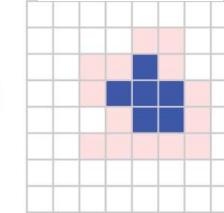


SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Erosion

- We put the element S on each pixel x of A
 - like convolution
- If **all pixels of S are onto object pixels (A)**, then the **central pixel** belongs to object
 - Otherwise (i.e. a mask pixel is background), set to background

Image A

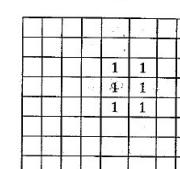
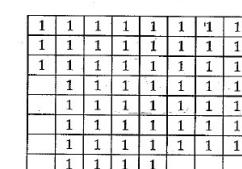
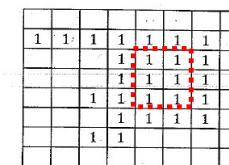
 $A \ominus S$ 

Structuring element S



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

2D example

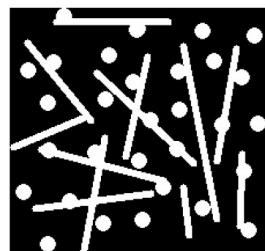


Shapiro & Stockman

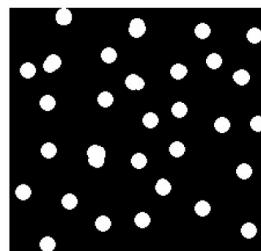


Opening

- Erode, then dilate
- Remove small objects, keep original shape



Before opening



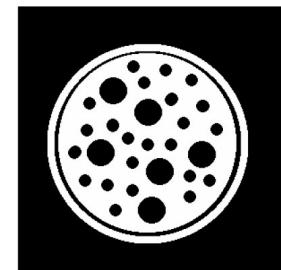
After opening



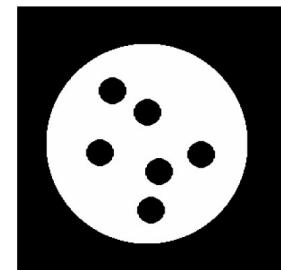
SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Closing

- Dilate, then erode
- Fill holes, but keep original shape



Before closing

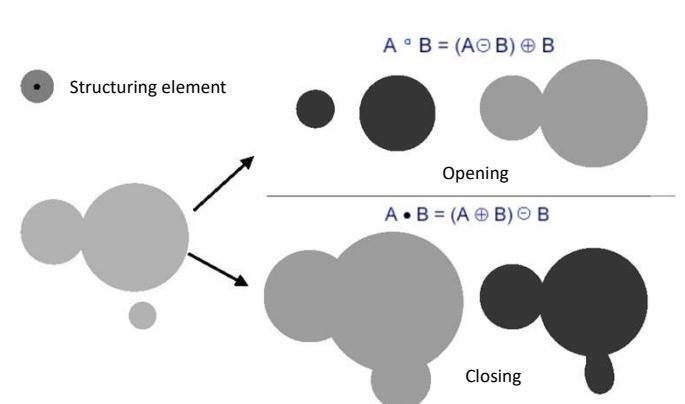


After closing



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Opening vs Closing

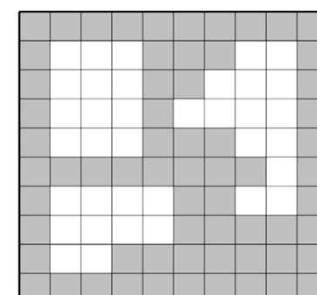


SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Connected component labeling

- We loop over all the image to give a **unique number (label)** for each region
- All pixels from the **same region** must have the **same number (label)**
- Objectifs:
 - Counting objects
 - Separating objects
 - Creating a mask for each object
 - ...

← *Background*
 ← *Segmented objects*

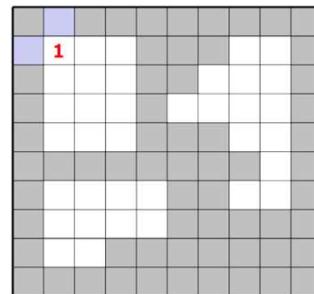
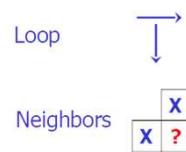


SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Connected component labeling

First loop over the image

- For each pixel in a region, we set
 - or the smallest label from its **top** or **left** neighbors
 - or a new label

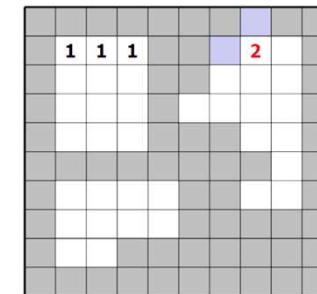
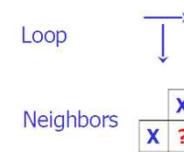


SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Connected component labeling

First loop over the image

- For each pixel in a region, we set
 - or the smallest label from its **top** or **left** neighbors
 - or a new label

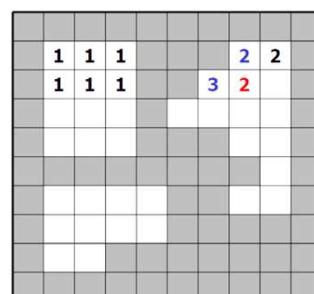
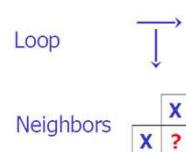


SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Connected component labeling

First loop over the image

- For each pixel in a region, we set
 - or the smallest label from its **top** or **left** neighbors
 - or a new label

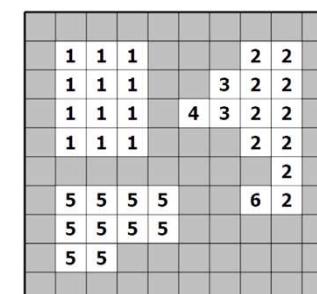
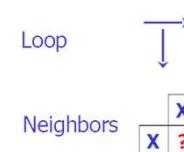


SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Connected component labeling

First loop over the image

- For each pixel in a region, we set
 - or the smallest label from its **top** or **left** neighbors
 - or a new label

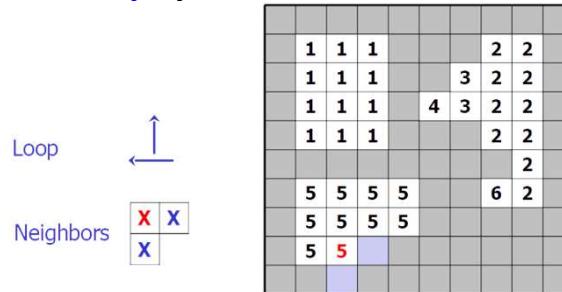


SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Connected component labeling

Second loop over the image

- For each pixel in a region, we set
 - the smallest from its *own label* and the labels from its *down* and *right* neighbors

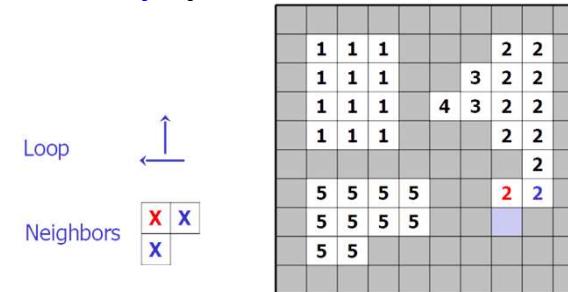


SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Connected component labeling

Second loop over the image

- For each pixel in a region, we set
 - the smallest from its *own label* and the labels from its *down* and *right* neighbors

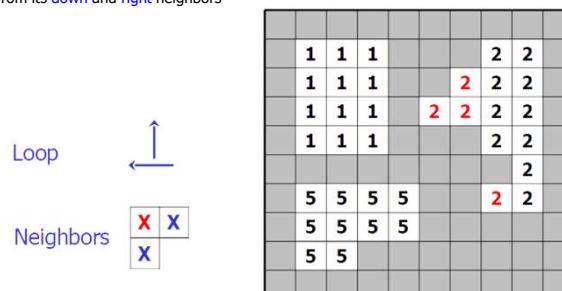


SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Connected component labeling

Second loop over the image

- For each pixel in a region, we set
 - the smallest from its *own label* and the labels from its *down* and *right* neighbors



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

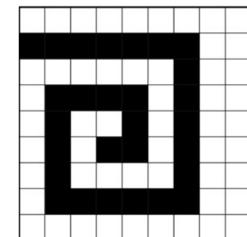
Connected component labeling

Two loops are enough?

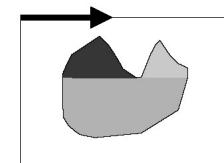
example: spiral region !

Solutions

- We continue, *go and back two ways*, until *no new change* in labels



- It is possible to do only one loop: manage a table of equivalences when 2 different labels are neighbors

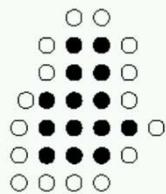
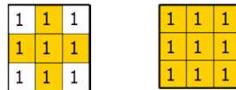


SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

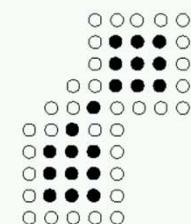
CC labeling: how many neighbors?

- Advice: Use different connexities for edges and regions

- 4-Connexity for regions
- 8-Connexity for edges



Region : 4-connected
Edge : 8-connected



Region : 8-connected
Edge : 4-connected

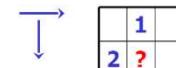


SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

CC labeling: how many neighbors?

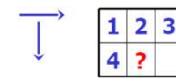
- Regions labeling

- We use 4-connexity
- Each loop, we compare 2 neighbors



- Edge labeling

- 8-connexity
- Each loop, we compare 4 neighbors



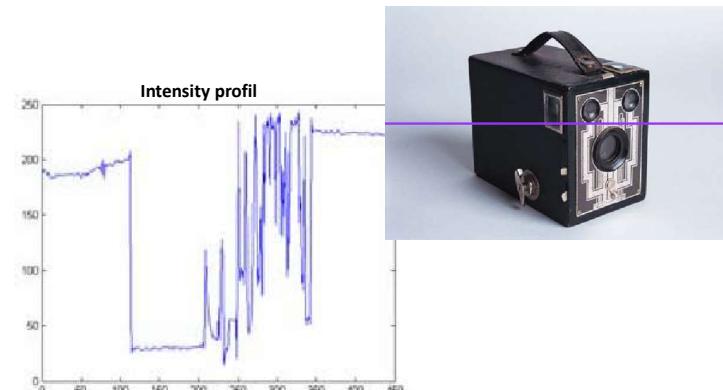
Content

- Rappel: digital image representation
- Point Processing
- Convolution and Linear filtering
- More neighborhood operators
- Image transforms
 - Frequency domain
 - Frequencies in images
 - Fourier transform
 - Frequency Processing (frequency filters)
 - PCA (additional reading)



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

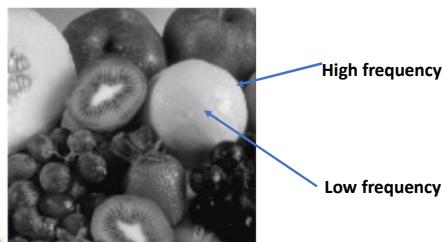
Frequencies in images



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Frequencies in images

- What are the (low/high) frequencies in an image?
 - Frequency = intensity change
 - Slow changes (homogeneous /blur regions): low frequency
 - fast/abrupt changes (edge, contour, noise): high frequency



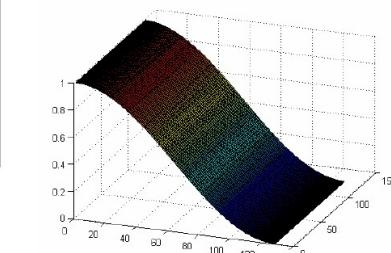
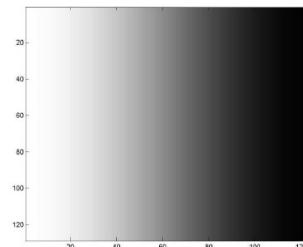
Most of energy concentrated in low frequencies



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

49

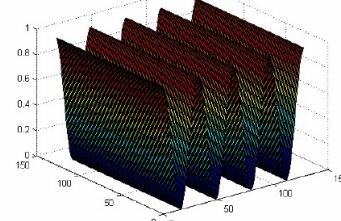
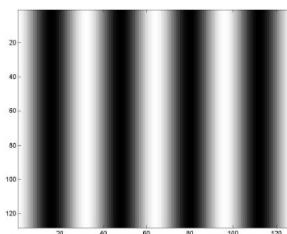
Low frequencies



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

50

High frequencies



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

51

Image spectral analysis

- An image is a visual signal
 - We can analyse the frequencies of the signal
- How?
 - we will create a new « image » which will contains all frequencies of the image
 - Like a 2D frequency graphic
 - The basic tool for it is the **Fourier Transform**
- We talk about the **frequency domain**, opposing to the **spatial domain** (image)



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

52

Frequencies in a signal

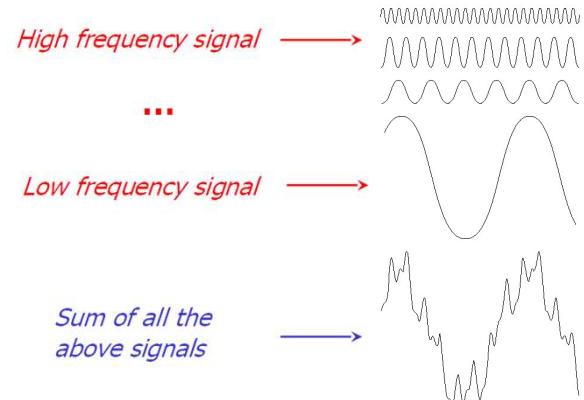


FIGURE 4.1 The function at the bottom is the sum of the four functions above it. Fourier's idea in 1807 that periodic functions could be represented as a weighted sum of sines and cosines was met with skepticism.



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Source : Gonzalez and Woods. Digital Image Processing. Prentice-Hall, 2002.

53

Fourier series

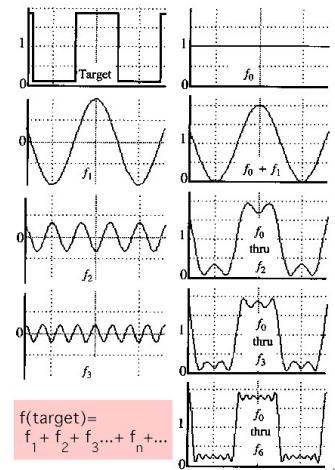
A bold idea (1807) - Jean Baptiste Joseph Fourier (1768-1830):

Any univariate function can be rewritten as a weighted sum of sines and cosines of different frequencies.

Our building block:

$$A \sin(\omega t) + B \cos(\omega t)$$

Add enough of them to get any signal $g(t)$ you want!



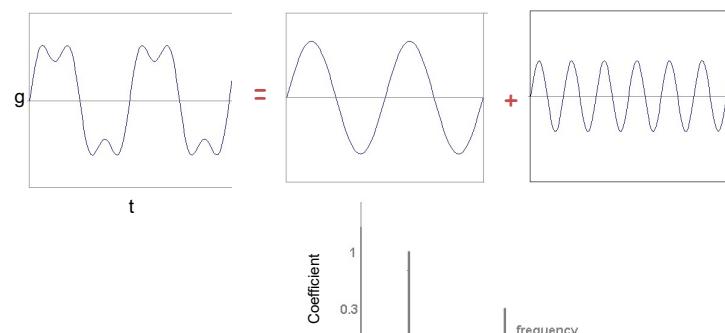
SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Hays

Example

$$t = [0, 2], f = 1$$

$$g(t) = \sin(2\pi f t) + (1/3)\sin(2\pi(3f)t)$$



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

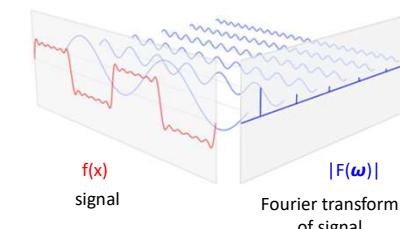
Slides: Efros

55

Fourier Transform

• Fourier transform is a mathematical transform that

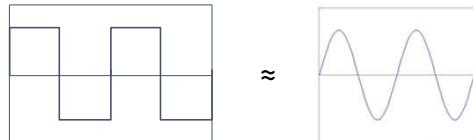
– Decomposes functions depending on space or time into functions depending on spatial or temporal frequency



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

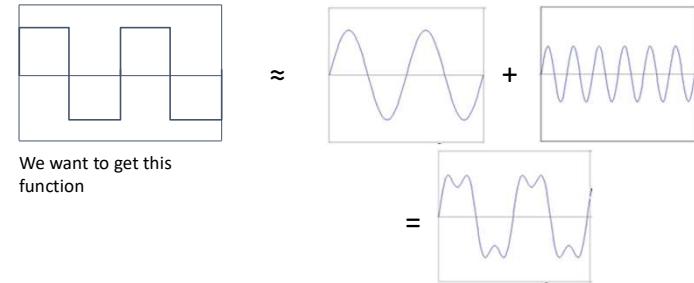
56

Fourier Series



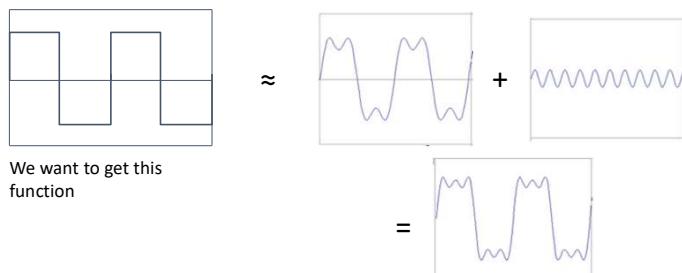
We want to get this function

Fourier Series



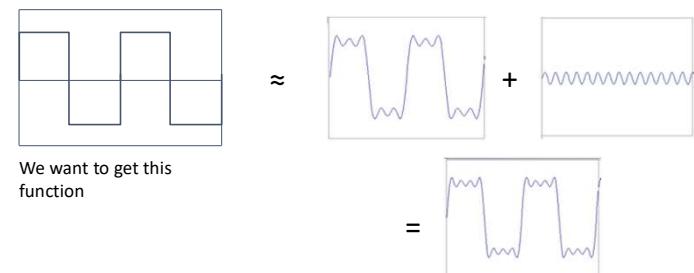
We want to get this function

Fourier Series



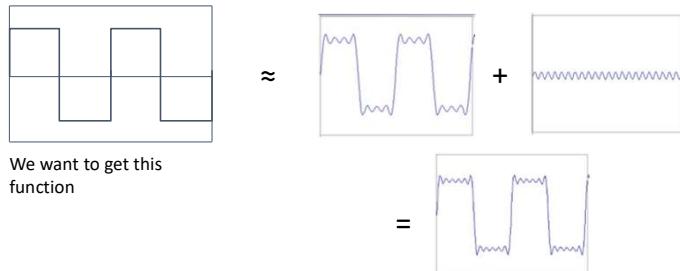
We want to get this function

Fourier Series

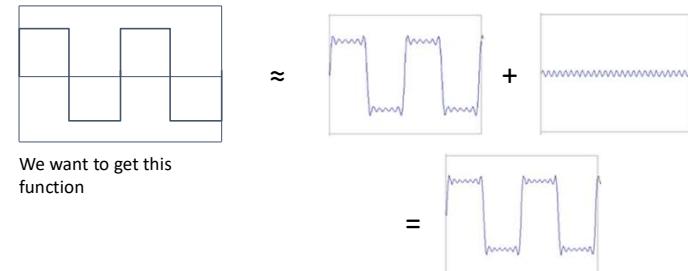


We want to get this function

Fourier Series



Fourier Series

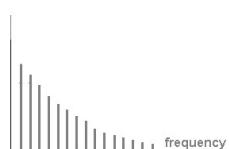


Fourier Series

$$= A \sum_{k=1}^{\infty} \frac{1}{k} \sin(2\pi kt)$$

We want to get this function

We'll get there in the limit



The math

$$\text{Fourier Transform : } F(\omega) = \int_{-\infty}^{+\infty} f(x) e^{-i\omega x} dx$$

$$\text{Inverse Fourier Transform : } f(x) = \frac{1}{2\pi} \int_{-\infty}^{+\infty} F(\omega) e^{i\omega x} d\omega$$

- Where are the sines and cosines? $e^{i\omega x} = \cos(\omega x) + i \sin(\omega x)$
- The result is a complex function $F(\omega) = R(\omega) + iI(\omega)$
- We've been showing only the **amplitude A (spectre)** so far:
- Phase is also encoded: $A = \pm \sqrt{R(\omega)^2 + I(\omega)^2}$

$$\phi = \tan^{-1} \frac{I(\omega)}{R(\omega)}$$

Magnitude and phase

- Fourier transform stores the **magnitude** and **phase** at each frequency
 - Magnitude encodes how much signal there is at a particular frequency
 - Phase encodes spatial information (indirectly)
 - For mathematical convenience, this is often notated in terms of real and complex numbers

Amplitude:

$$A = \pm \sqrt{R(\omega)^2 + I(\omega)^2}$$

Phase:

$$\phi = \tan^{-1} \frac{I(\omega)}{R(\omega)}$$



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Slide by Rober Pless

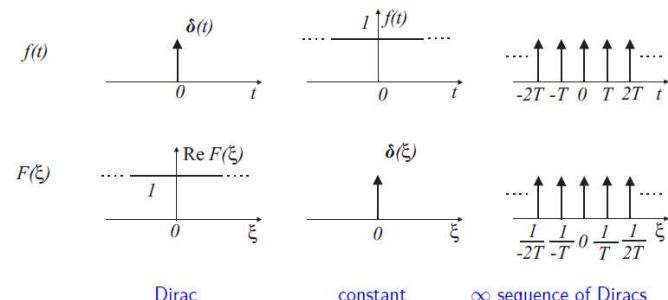
Discrete Fourier transform

$$H_{f_j} = \frac{1}{N} \sum_k h_{t_k} e^{2\pi i f_j t_k}$$

$$h_{t_j} = \frac{1}{N} \sum_k H_{f_k} e^{-2\pi i f_k t_j}$$

where the t_k are the time corresponding to my signal in the time domain h_{t_k} , f_k are the corresponding frequency to my signal in the frequency domain, and N is the number of points of the signal data.

Basic Fourier Transform pairs

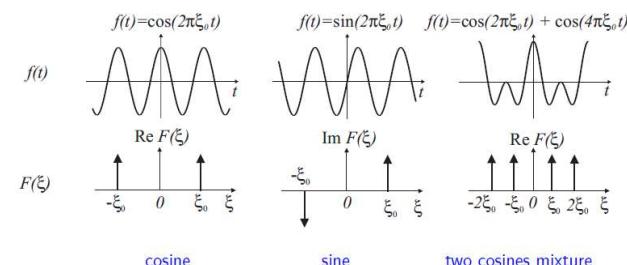


SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

67

Source: Václav Hlaváč - Fourier transform, in 1D and in 2D

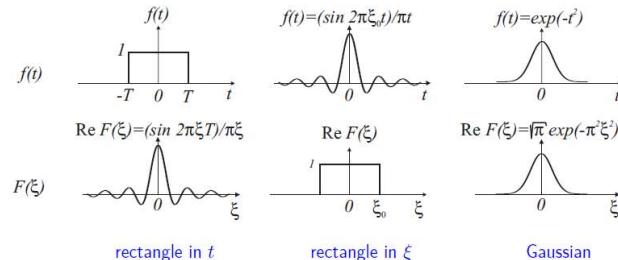
Basic Fourier Transform pairs



Source: Václav Hlaváč - Fourier transform, in 1D and in 2D

68

Basic Fourier Transform pairs



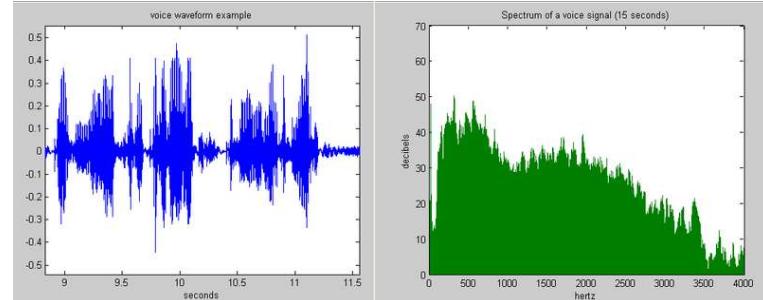
SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Source: Václav Hlaváč - Fourier transform, in 1D and in 2D

69

Example: Music

- We think of music in terms of frequencies at different magnitudes



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Slide: Hoiem

2D FFT

- Continuous FFT:

$$F(u, v) = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} f(x, y) e^{-2\pi i(xu+yv)} dx dy$$

- Inverse FFT:

$$f(x, y) = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} F(u, v) e^{2\pi i(xu+yv)} du dv$$



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

2D FFT - discrete

Direct transform

$$F(u, v) = \frac{1}{MN} \sum_{m=0}^{M-1} \sum_{n=0}^{N-1} f(m, n) \exp \left[-2\pi i \left(\frac{mu}{M} + \frac{nv}{N} \right) \right],$$

$$u = 0, 1, \dots, M-1, \quad v = 0, 1, \dots, N-1,$$

Inverse transform

$$f(m, n) = \sum_{u=0}^{M-1} \sum_{v=0}^{N-1} F(u, v) \exp \left[2\pi i \left(\frac{mu}{M} + \frac{nv}{N} \right) \right],$$

$$m = 0, 1, \dots, M-1, \quad n = 0, 1, \dots, N-1.$$



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Image Fourier transform

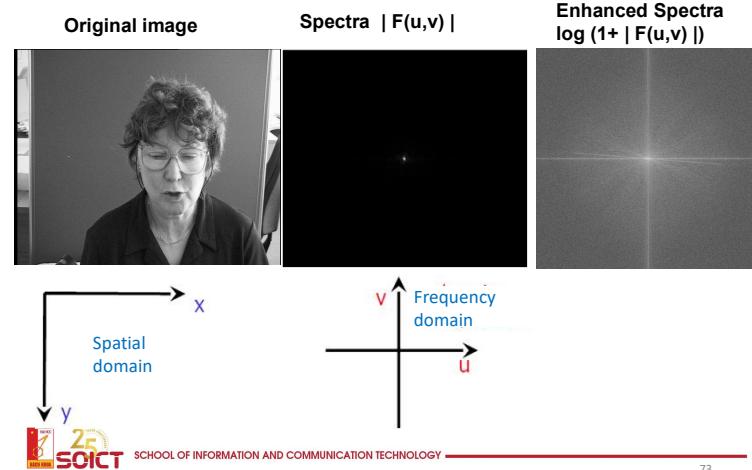


Image Fourier transform

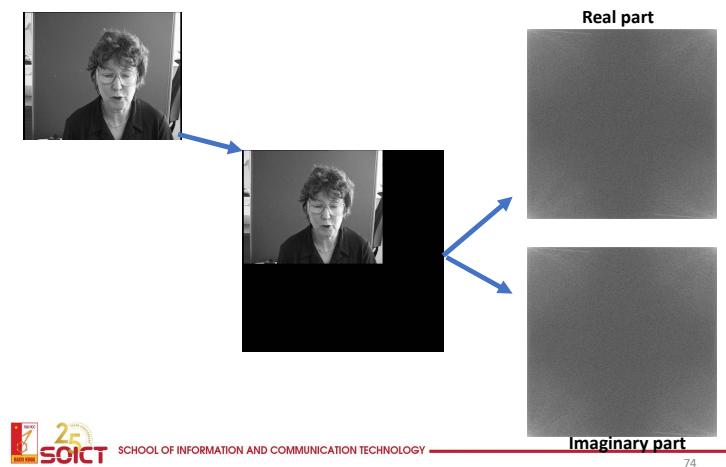
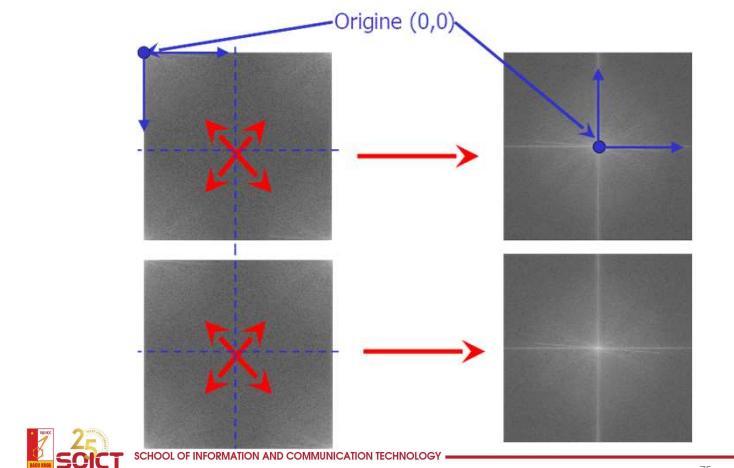
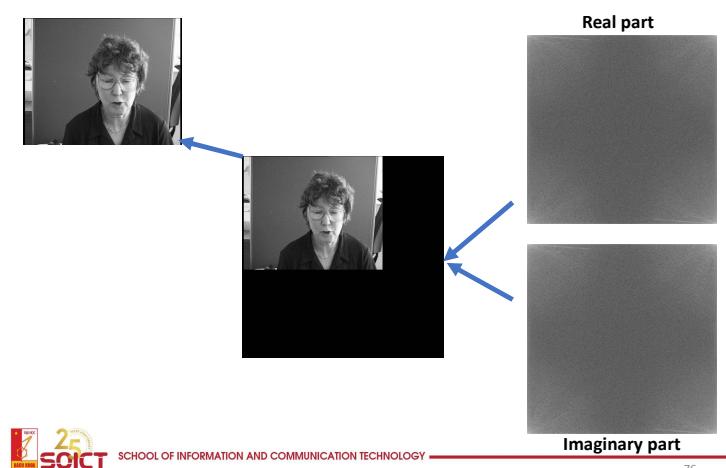


Image Fourier transform

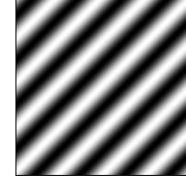
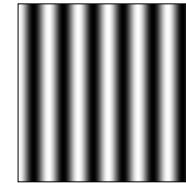
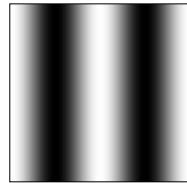


Inverse Fourier transform

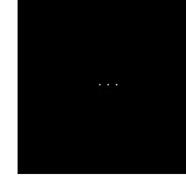
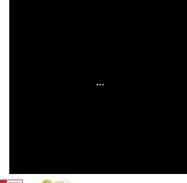


Fourier analysis in images

Intensity images (spatial domain)



Fourier images (spectral images – amplitude images)

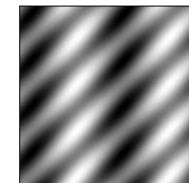
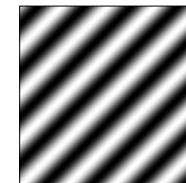
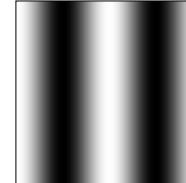


<http://sharp.bu.edu/~slehar/fourier/fourier.html#filtering> More: <http://www.cs.unm.edu/~brayer/vision/fourier.html>

77

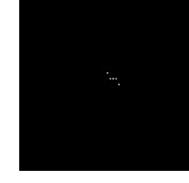
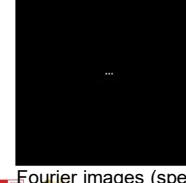
Signals can be composed

Intensity images (spatial domain)



+

=



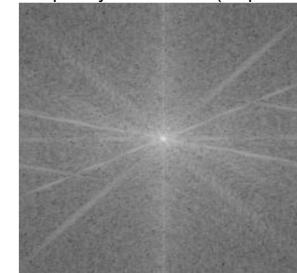
<http://sharp.bu.edu/~slehar/fourier/fourier.html#filtering>

<http://sharp.bu.edu/~slehar/fourier/fourier.html#filtering>

78

Fourier Transform of an image

Natural image

 $f(x,y)$ Fourier decomposition
Frequency coefficients (amplitude) $|F(\omega)|$

What does it mean to be at pixel x,y ?

What does it mean to be more or less bright in the Fourier decomposition image?

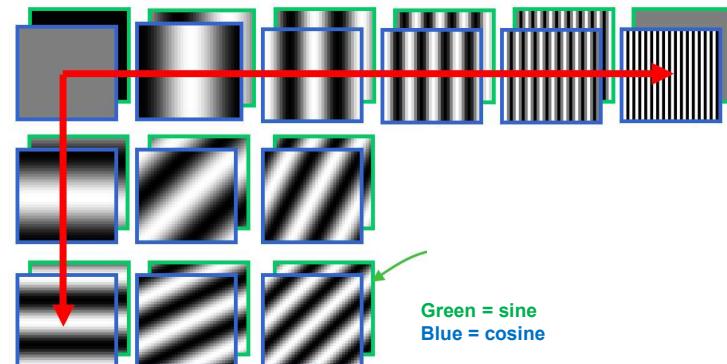
Slide by Steve Seitz



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Fourier Bases

Teases away ‘fast vs. slow’ changes in the image.

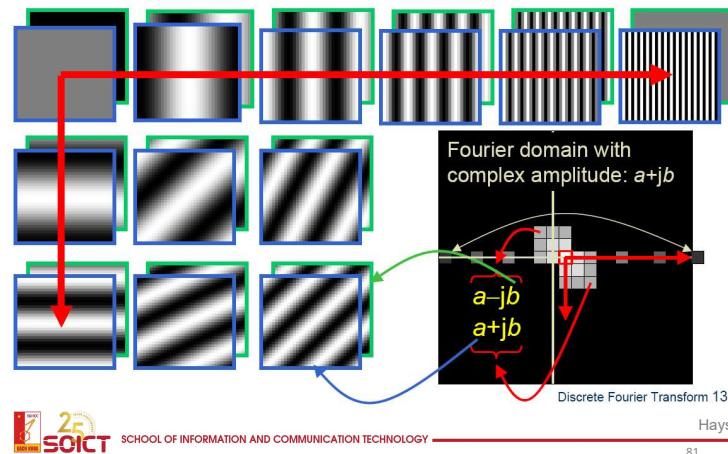


This change of basis is the Fourier Transform

Hays

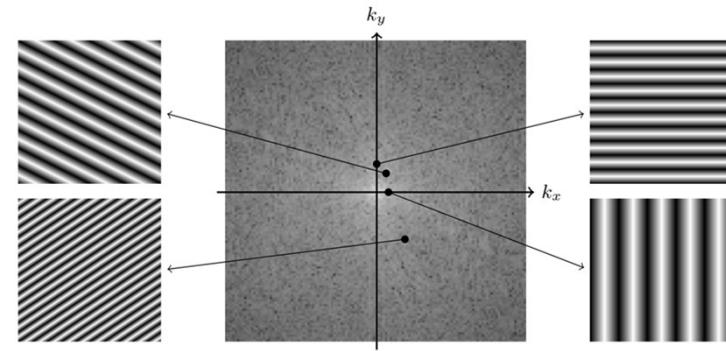
80

Fourier Bases



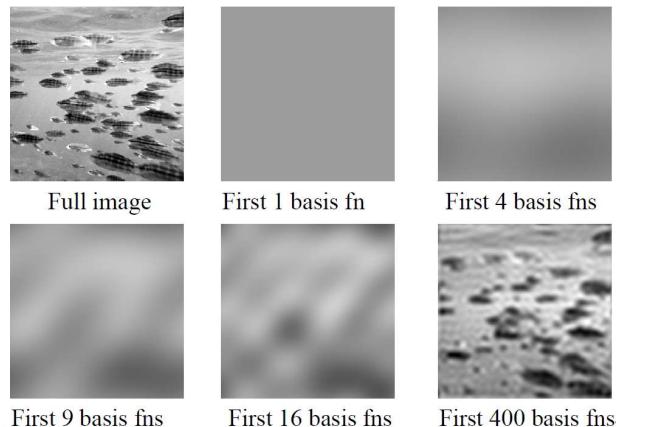
SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

2D Fourier Transform



Slide by Steve Seitz

Basis reconstruction

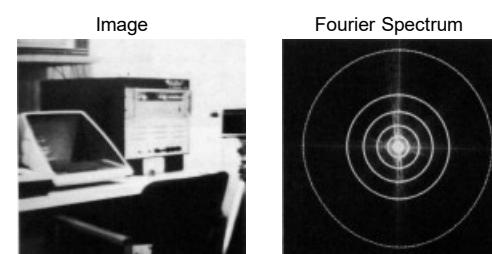


Danny Alexander



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

2D Fourier transform



Percentage of image power enclosed in circles (small to large) :
90, 95, 98, 99, 99.5, 99.9

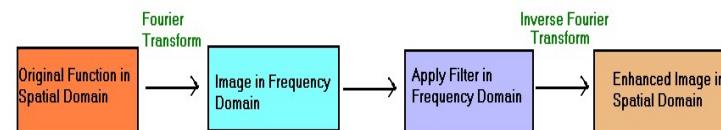
Most of energy concentrated in low frequencies

Danny Alexander

84

Image filtering in the frequency domain

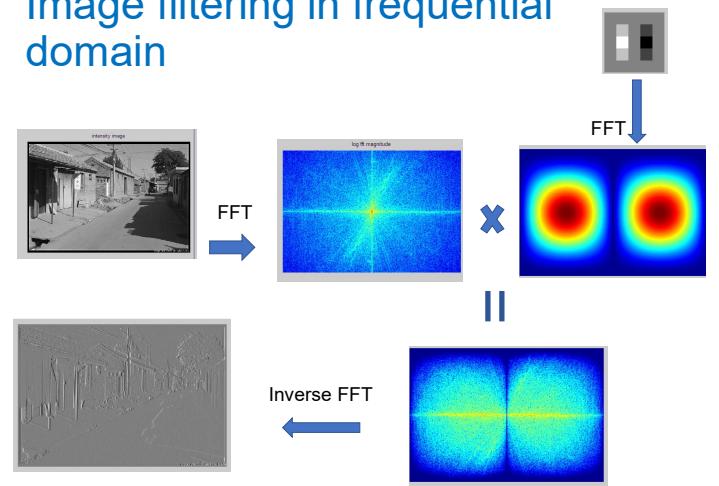
- We will be dealing only with functions (images) of finite duration so we will be interested only in Fourier Transform



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

85

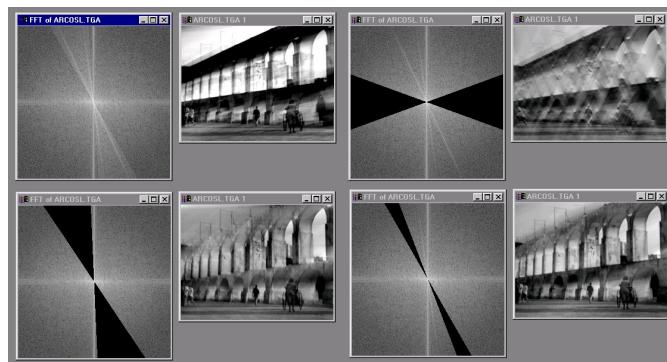
Image filtering in frequency domain



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Slide by Derek Hoiem

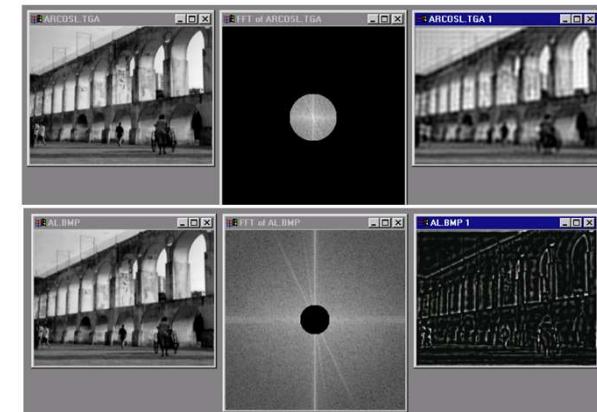
Now we can edit frequencies!



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

87

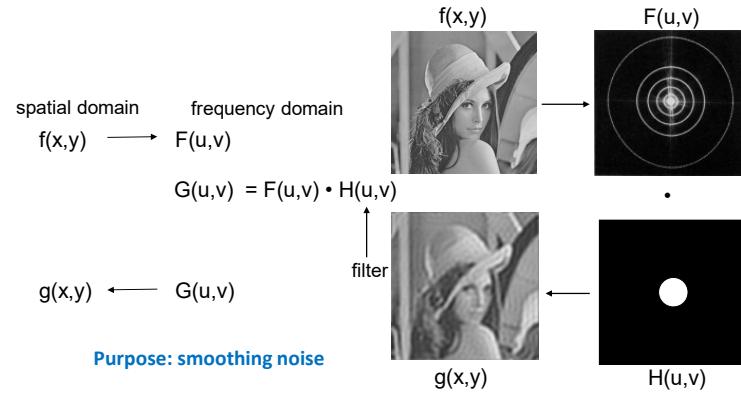
Low-pass and high-pass filtering



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Slide by Alexei Efros

Low-pass filter



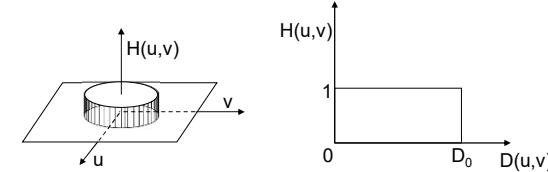
SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

89

$H(u,v)$ - Ideal low-pass filter

$$H(u,v) = \begin{cases} 1 & D(u,v) \leq D_0 \\ 0 & D(u,v) > D_0 \end{cases} \quad D(u,v) = \sqrt{u^2 + v^2}$$

D_0 = cut off frequency



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

90

Blurring - Ideal low-pass filters

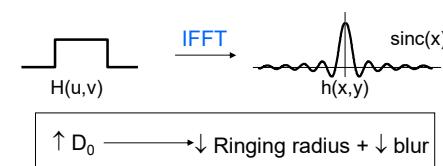


SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

91

The ringing problem

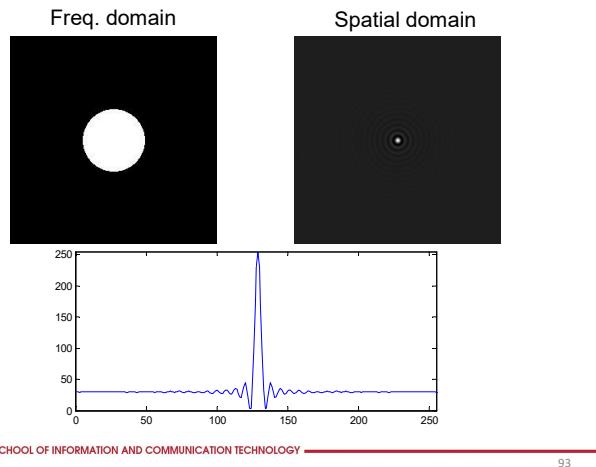
$$\begin{aligned} G(u,v) &= F(u,v) \cdot H(u,v) \\ &\downarrow \text{Convolution Theorem} \\ g(x,y) &= f(x,y) * h(x,y) \end{aligned}$$



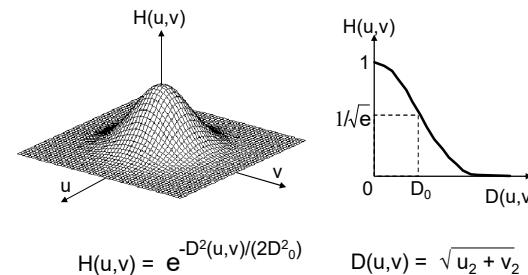
SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

92

The ringing problem

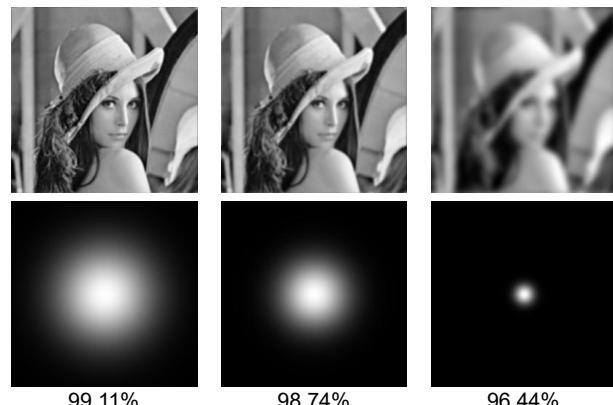


$H(u,v)$ - Gaussian filter

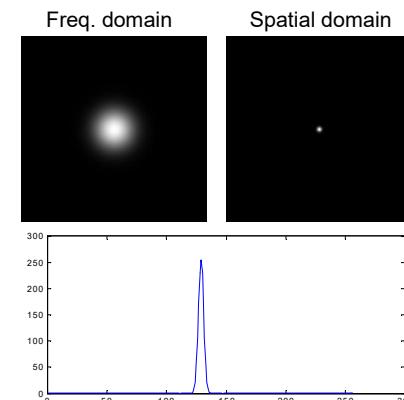


Softer Blurring + no Ringing

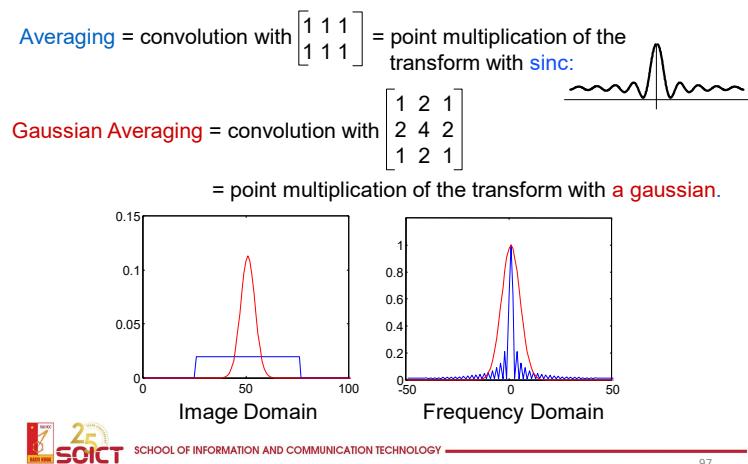
Blurring - Gaussian lowpass filter



The Gaussian lowpass filter



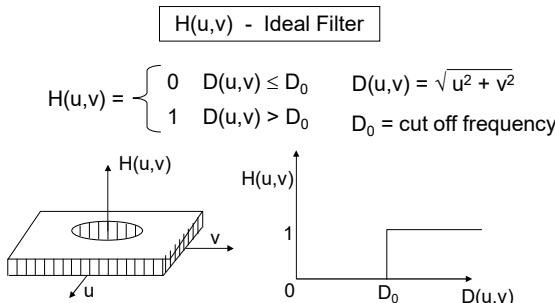
Blurring in the Spatial Domain



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

97

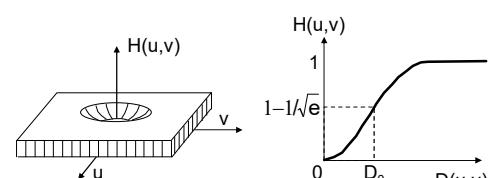
High-pass filter



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

98

High-pass gaussian filter



$$H(u,v) = 1 - e^{-D^2(u,v)/(2D_0^2)}$$

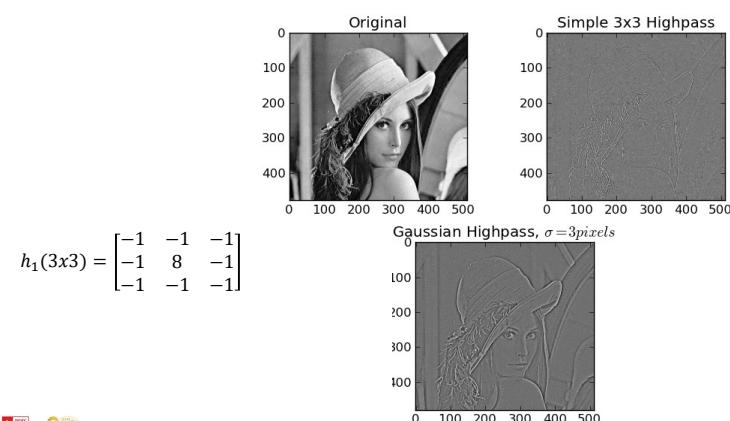
$$D(u,v) = \sqrt{u^2 + v^2}$$



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

99

High-pass filtering



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

100

High pass filtering

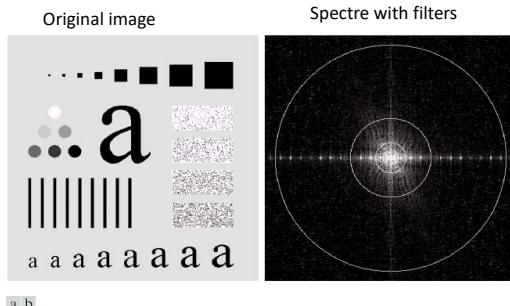


FIGURE 4.11 (a) An image of size 500×500 pixels and (b) its Fourier spectrum. The superimposed circles have radii values of 5, 15, 30, 80, and 230, which enclose 92.0, 94.6, 96.4, 98.0, and 99.5% of the image power, respectively.

Source : Gonzalez and Woods. *Digital Image Processing*. Prentice-Hall, 2002.



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

101

High pass filtering

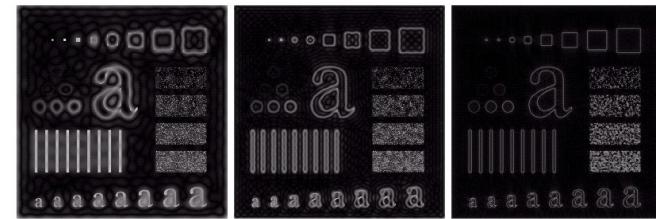


FIGURE 4.24 Results of ideal highpass filtering the image in Fig. 4.11(a) with $D_0 = 15, 30$, and 80 , respectively. Problems with ringing are quite evident in (a) and (b).

Source : Gonzalez and Woods. *Digital Image Processing*. Prentice-Hall, 2002.



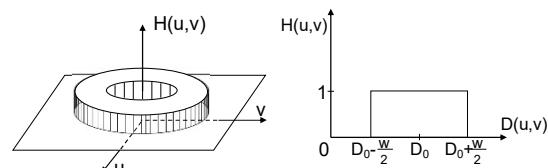
SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

102

Band-pass filtering

$$H(u,v) = \begin{cases} 0 & D(u,v) \leq D_0 - \frac{w}{2} \\ 1 & D_0 - \frac{w}{2} \leq D(u,v) \leq D_0 + \frac{w}{2} \\ 0 & D(u,v) > D_0 + \frac{w}{2} \end{cases} \quad D(u,v) = \sqrt{u^2 + v^2}$$

D_0 = cut off frequency
 w = band-width



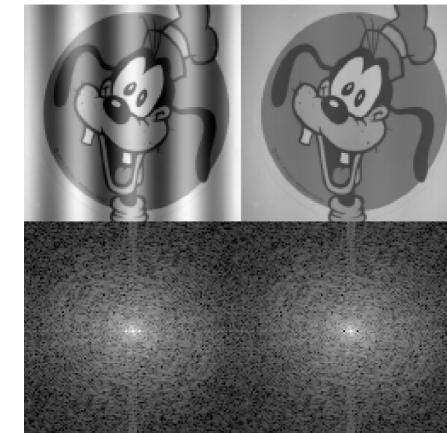
Can be obtained by multiplying the filter functions of a low-pass and of a high-pass in the frequency domain



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

103

Removing sinus noise

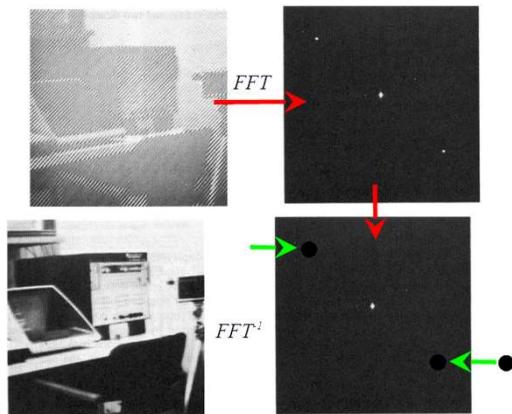


Brayer



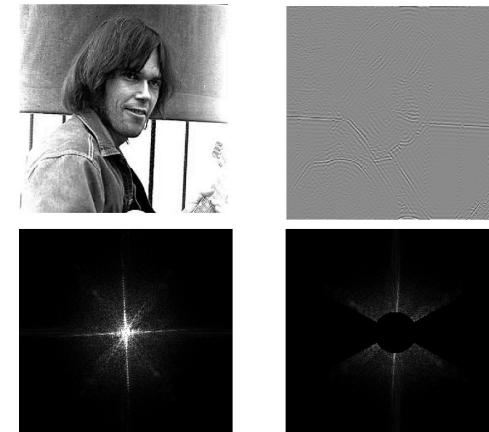
SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Removing sinus noise



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

High-pass filtering + orientation



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Hybrid Images

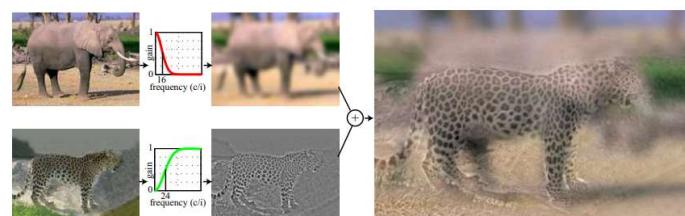


Figure 2: hybrid images are generated by superimposing two images at two different spatial scales: the low-spatial scale is obtained by filtering one image with a low-pass filter, and the high spatial scale is obtained by filtering a second image with a high-pass filter. The final hybrid image is composed by adding these two filtered images.

A. Oliva, A. Torralba, P.G. Schyns, SIGGRAPH 2006



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Content

- Rappel: digital image representation
- Point Processing
- Convolution and Linear filtering
- More neighborhood operators
- Image transforms
 - Frequency domain
 - PCA (additional reading)
 - PCA
 - Example of using PCA for face recognition



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Principle Component Analysis - PCA (Karhunen-Loeve transformation)

- PCA transforms the original input space into a lower dimensional space
 - By constructing dimensions that are linear combinations of the given features
- The objective: consider independent dimensions along which data have largest variance (i.e., greatest variability)



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Slide by Jana Kosecka

Principal Component Analysis (cont.)

- PCA enables transform a number of possibly correlated variables into a smaller number of uncorrelated variables called principal components
- The first principal component accounts for as much of the variability in the data as possible
- Each succeeding component (orthogonal to the previous ones) accounts for as much of the remaining variability as possible



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Slide by Jana Kosecka

Principal Component Analysis (cont.)

- PCA is the most commonly used dimension reduction technique.
- Data samples x_1, \dots, x_N
- Compute the mean $\bar{x} = \frac{1}{N} \sum_{i=1}^N x_i$
- Compute the covariance matrix:

$$\Sigma_x = \frac{1}{N} \sum_{i=1}^N (x_i - \bar{x})(x_i - \bar{x})^T$$



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Slide by Jana Kosecka

Principal Component Analysis (cont.)

- Compute the eigenvalues λ and eigenvectors e of the matrix Σ_x
- Solve $\Sigma_x e = \lambda e$
- Order them by magnitude: $\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_N$.
- PCA reduces the dimension by keeping direction e such that $\lambda < T$.



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

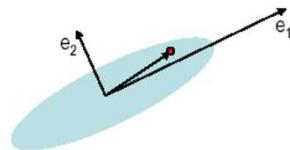
Slide by Jana Kosecka

Principal Component Analysis (cont.)

- For many datasets, most of the eigenvalues are negligible and can be discarded.

The eigenvalue λ measures the variation in the direction of corresponding eigenvector

Example:
 $\lambda_1 \neq 0, \lambda_2 = 0$.



Slide by Jana Kosecka



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Principal Component Analysis (cont.)

- How to get uncorrelated components which capture most of the variance
- Project the data onto the selected eigenvectors:
 $y_i = e_i^T (x_i - \bar{x})$
- If we consider first M eigenvectors we get new lower dimensional representation
 $[y_1, \dots, y_M]$
- Proportion covered by first M eigenvalues

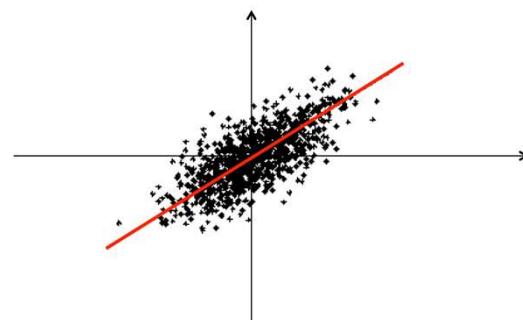
$$\frac{\sum_{i=1}^M \lambda_i}{\sum_{i=1}^N \lambda_i}$$

Slide by Jana Kosecka



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Illustration of PCA



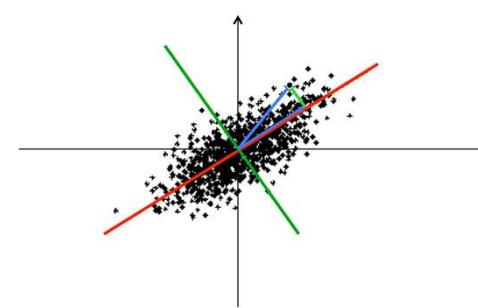
First principal component of a two-dimensional data set.

Slide by Jana Kosecka



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Illustration of PCA



Second principal component of a two-dimensional data set.

Slide by Jana Kosecka



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Determining the number of components

- Plot the eigenvalues
 - each eigenvalue is related to the amount of variation explained by the corresponding axis (eigenvector)
 - If the points on the graph tend to level out (show an “elbow” shape), these eigenvalues are **usually close enough to zero that they can be ignored**

Content

- Rappel: digital image representation
- Point Processing
- Convolution and Linear filtering
- More neighborhood operators
- Image transforms
 - Frequency domain
 - PCA (additional reading)
 - PCA
 - Example of using PCA for face recognition

The space of all face images

- When viewed as vectors of pixel values, face images are extremely high-dimensional
 - 100×100 image = 10,000 dimensions
- However, relatively few 10,000-dimensional vectors correspond to valid face images
- We want to effectively model the subspace of face images



Eigenfaces: Key idea

- Assume that most face images lie on a low-dimensional subspace determined by the first k ($k < d$) directions of maximum variance
- Use PCA to determine the vectors or “eigenfaces” u_1, \dots, u_k that span that subspace
- Represent all face images in the dataset as linear combinations of eigenfaces

Eigenfaces example- Traning images



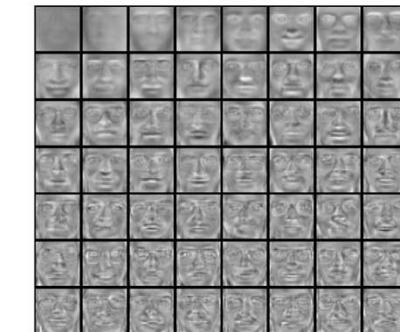
Slide by Jana Kosecka



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Eigenfaces example

Top eigenvectors: u_1, \dots, u_k

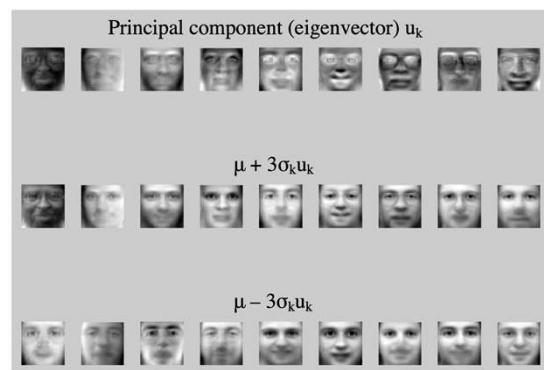


Slide by Jana Kosecka



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Eigenfaces example



Slide by Jana Kosecka



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Eigenfaces examples

- Representation



$$(w_{i1}, \dots, w_{ik}) = (u_1^T(x_i - \mu), \dots, u_k^T(x_i - \mu))$$

- Reconstruction

$$\begin{aligned} \hat{x} &= \mu + [w_1 u_1 + w_2 u_2 + w_3 u_3 + w_4 u_4 + \dots] \\ &= \mu + w_1 u_1 + w_2 u_2 + w_3 u_3 + w_4 u_4 + \dots \end{aligned}$$

Slide by Jana Kosecka



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Recognition with eigenfaces

- Process labeled training images:
- Find mean μ and covariance matrix Σ
- Find k principal components (eigenvectors of Σ) u_1, \dots, u_k
- Project each training image x_i onto subspace spanned by principal components:
 $(w_{i1}, \dots, w_{ik}) = (u_1^T(x_i - \mu), \dots, u_k^T(x_i - \mu))$
- Given novel image x :
- Project onto subspace:
 $(w_1, \dots, w_k) = (u_1^T(x - \mu), \dots, u_k^T(x - \mu))$
- Optional: check reconstruction error $x - \hat{x}$ to determine whether image is really a face
- Classify as closest training face in k -dimensional subspace

